



Distributed Systems

Lab 2 Report

Deploying a request splitting ambassador and a load balancer with Kubernetes

Hemshikha Sultoo

100670616

Part 1

The problems being solved are:

- 1) Multiple disparate services
- 2) Multiple instances of the same service
- 3) Multiple versions of the same service

Client applications only need to know about a single endpoint and communicate with a single endpoint with this pattern.

Requirements:

- 1) Client needs to consume multiple services that can be accessed behind a gateway
 - 2) Use single endpoint to simplify client applications/
 - 3) Route requests from externally addressable endpoints to internal virtual endpoints
 - 4) Clients require services running in many regions for availability or latency benefits.
 - 5) Clients consume a variable number of service instances.
 - 6) Create a deployment strategy so that clients can access many versions of the service simultaneously.
-

Part 2:

Creating the web-deployment.yaml file

Creating the deployment on the GKE

Checking pods

```
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl create -f web-deployment.yaml
deployment.apps/web-deployment created
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
mongo-0                             1/1     Running   0           30h
mongo-deployment-0                  1/1     Running   0           30h
mongodb-0                           1/1     Running   0           30h
web-deployment-6fdbb5c6bb-7x2mv     1/1     Running   0           24s
web-deployment-6fdbb5c6bb-k87j6     1/1     Running   0           24s
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$
```

Create clusterIP

```
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl expose deployment web-deployment --port=80 --type=ClusterIP --name web-deployment
service/web-deployment exposed
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$
```

Creating a deployment for experiment-deployment.yaml.

Checking pods

```
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl create -f experiment-deployment.yaml
deployment.apps/experiment-deployment created
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$
```

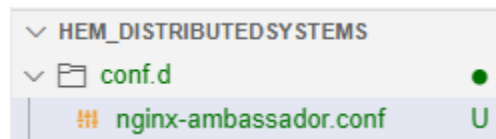
```
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
experiment-deployment-7b47cbd668-9ljdm 1/1     Running   0           69s
experiment-deployment-7b47cbd668-r2rr9 1/1     Running   0           69s
mongo-0                                1/1     Running   0           30h
mongo-deployment-0                     1/1     Running   0           30h
mongodb-0                              1/1     Running   0           30h
web-deployment-6fdbb5c6bb-7x2mv        1/1     Running   0           6m37s
web-deployment-6fdbb5c6bb-k87j6        1/1     Running   0           6m37s
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$
```

Create clusterIP again

```
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl expose deployment experiment-deployment --port=80
--type=ClusterIP --name experiment-deployment
service/experiment-deployment exposed
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$
```

Create request splitter

Part a:



Part b:

```
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl create configmap ambassador-config --from-file=conf.d
configmap/ambassador-config created
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$
```

Part c and d:

```
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl create -f ambassador-deployment.yaml
deployment.apps/ambassador-deployment created
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$
```

Part e:

```
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl expose deployment ambassador-deployment --port=80
--type=LoadBalancer
service/ambassador-deployment exposed
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$
```

Part f:

```
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
ambassador-deployment-66db4f7766-8c9vz 1/1     Running   0           7m58s
ambassador-deployment-66db4f7766-drh5k 1/1     Running   0           7m58s
experiment-deployment-7b47cbd668-9ljdm 1/1     Running   0           24m
experiment-deployment-7b47cbd668-r2rr9 1/1     Running   0           24m
mongo-0                                1/1     Running   0           30h
mongo-deployment-0                     1/1     Running   0           30h
mongodb-0                               1/1     Running   0           30h
web-deployment-6fdbb5c6bb-7x2mv        1/1     Running   0           29m
web-deployment-6fdbb5c6bb-k87j6        1/1     Running   0           29m
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl get deployments
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
ambassador-deployment 2/2     2             2           8m10s
experiment-deployment 2/2     2             2           24m
web-deployment        2/2     2             2           29m
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl get services
NAME                TYPE          CLUSTER-IP   EXTERNAL-IP   PORT(S)          AGE
ambassador-deployment LoadBalancer  10.80.4.204   35.203.83.103 80:31429/TCP      86s
experiment-deployment ClusterIP      10.80.14.244  <none>         80/TCP            20m
kubernetes           ClusterIP      10.80.0.1     <none>         443/TCP           13d
mongo-nodeport-svc   NodePort      10.80.12.236  <none>         27017:32000/TCP   13d
mongo-service        LoadBalancer  10.80.13.8    35.234.254.180 27017:31794/TCP   13d
mongodb              ClusterIP      None          <none>         27017/TCP         13d
mongodb-service      LoadBalancer  10.80.12.59   34.152.21.61   27017:30683/TCP   13d
web-deployment       ClusterIP      10.80.10.46   <none>         80/TCP            26m
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$
```

Step 5

Part a:

```
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl get services
NAME                TYPE          CLUSTER-IP   EXTERNAL-IP   PORT(S)          AGE
ambassador-deployment LoadBalancer  10.80.4.204   35.203.83.103 80:31429/TCP      86s
experiment-deployment ClusterIP      10.80.14.244  <none>         80/TCP            20m
kubernetes           ClusterIP      10.80.0.1     <none>         443/TCP           13d
mongo-nodeport-svc   NodePort      10.80.12.236  <none>         27017:32000/TCP   13d
mongo-service        LoadBalancer  10.80.13.8    35.234.254.180 27017:31794/TCP   13d
mongodb              ClusterIP      None          <none>         27017/TCP         13d
mongodb-service      LoadBalancer  10.80.12.59   34.152.21.61   27017:30683/TCP   13d
web-deployment       ClusterIP      10.80.10.46   <none>         80/TCP            26m
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$
```

Part b:

```

hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ curl http://35.203.83.103
<html>
<head>
  <title>Welcome to Azure Container Instances!</title>
</head>
<style>
h1 {
  color: darkblue;
  font-family:arial, sans-serif;
  font-weight: lighter;
}
</style>

<body>

<div align="center">
<h1>Welcome to Azure Container Instances!</h1>

<svg id="Layer_1" data-name="Layer 1" xmlns="http://www.w3.org/2000/svg" viewBox="0 0 49.8 49.9" width="250px" height="250px">
  <title>ContainerInstances_rgb_UI</title>
  <path d="M41.9,11.368A11.929,11.929,0,0,0,20.3,5.061a9.444,9.444,0,0,0-14.932,9.8A8.969,8.969,0,0,0,9.064,32H39.442A10.463,10.463,0,0,0,41.9,11.368Z" transform="translate(-0.1 -0.1)" fill="#fff"/>
  <path d="M41.9,11.368A11.929,11.929,0,0,0,20.3,5.061a9.444,9.444,0,0,0-14.932,9.8A8.969,8.969,0,0,0,9.064,32H39.442A10.463,10.463,0,0,0,41.9,11.368Z" transform="translate(-0.1 -0.1)" fill="#27a9e1" opacity="0.6" style="isolation:isolate"/>
  <path d="M13,22a1,1,0,0,0-1,1V49a1,1,0,0,1,1H37a1,1,0,0,1-1V23a1,1,0,0,0-1-1Z" transform="translate(-0.1 -0.1)" fill="#672a7a"/>
  <path d="M26.95,16" transform="translate(-0.1 -0.1)" fill="none"/>
  <path d="M34.95,20" transform="translate(-0.1 -0.1)" fill="none"/>
  <polygon points="22.9 21.9 22.9 14.9 19.9 14.9 24.9 7.9 29.9 14.9 26.9 14.9 26.9 21.9 22.9 21.9" fill="#fff"/>
  <path d="M26.95,16" transform="translate(-0.1 -0.1)" fill="#814a98"/>
  <path d="M33,25H15V47H35V25ZM21,45H17V27h4Zm6,0H23V27h4Zm6,0H29V27h4Z" transform="translate(-0.1 -0.1)" fill="#b92025" opacity="0.3" style="isolation:isolate"/>
  <path d="M33,25H15V47H35V25ZM21,45H17V27h4Zm6,0H23V27h4Zm6,0H29V27h4Z" transform="translate(-0.1 -0.1)" fill="#fff" style="isolation:isolate"/>
</svg>
</div>

</body>

```

Part c:

```

hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ for _ in {1..20}; do curl http://35.203.83.103 -s > output.txt; done

hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ 

```

Part d:

Part 3

Step 1

Part a and b:

```
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl create -f loadbalancer-deployment.yaml
deployment.apps/loadbalancer-deployment created
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$
```

Part c

```
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl get pods --output=wide
NAME                                READY   STATUS    RESTARTS   AGE   IP              NODE                                NOMINATED NODE   READINESS GATES
loadbalancer-deployment-6676f9ccf6-7pws6   1/1     Running   0           61s   10.76.1.8       gke-openfaas-default-pool-392835dd-ryhh   <none>           <none>
loadbalancer-deployment-6676f9ccf6-gh27v   1/1     Running   0           61s   10.76.2.11      gke-openfaas-default-pool-392835dd-rjso   <none>           <none>
loadbalancer-deployment-6676f9ccf6-s4l6a   1/1     Running   0           60s   10.76.3.9       gke-openfaas-default-pool-392835dd-mx0t   <none>           <none>
mongo-0                                    1/1     Running   0           31h   10.76.3.6       gke-openfaas-default-pool-392835dd-mx0t   <none>           <none>
mongo-deployment-0                        1/1     Running   0           31h   10.76.1.4       gke-openfaas-default-pool-392835dd-ryhh   <none>           <none>
mongodb-0                                 1/1     Running   0           31h   10.76.3.7       gke-openfaas-default-pool-392835dd-mx0t   <none>           <none>
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$
```

Step 2

Part a and b

```
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl expose deployment loadbalancer-deployment --port=8080 --type=LoadBalancer
service/loadbalancer-deployment exposed
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl get services --watch
NAME                                TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
kubernetes                          ClusterIP    10.80.0.1        <none>            443/TCP          13d
loadbalancer-deployment             LoadBalancer 10.80.11.182     <pending>        8080:31110/TCP   19s
mongo-nodeport-svc                 NodePort     10.80.12.236     <none>            27017:32000/TCP  13d
mongo-service                      LoadBalancer 10.80.13.8       35.234.254.180   27017:31794/TCP  13d
mongodb                            ClusterIP    None             <none>            27017/TCP        13d
mongodb-service                   LoadBalancer 10.80.12.59      34.152.21.61     27017:30683/TCP  13d
loadbalancer-deployment             LoadBalancer 10.80.11.182     35.203.83.103    8080:31110/TCP   45s
^Chem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$
```

Step 3:

```
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ curl http://35.203.83.103:8080/dog
A quadruped of the genus Canis, esp. the domestic dog (C.familiaris).hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ curl http://35.203.83.103:8080/storey
See Story.hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ curl http://35.203.83.103:8080/storey
See Story.hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ curl http://35.203.83.103/storey
```

Step 4

```
See Story.hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl delete deployment loadbalancer-deployment
deployment.apps "loadbalancer-deployment" deleted
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl delete service loadbalancer-deployment
service "loadbalancer-deployment" deleted
```

Design - Autoscaling

```
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl get hpa
NAME          REFERENCE                TARGETS  MINPODS  MAXPODS  REPLICAS  AGE
php-apache    Deployment/php-apache     0%/50%   1         10        1          12m
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl get hpa php-apache --watch
NAME          REFERENCE                TARGETS  MINPODS  MAXPODS  REPLICAS  AGE
php-apache    Deployment/php-apache     43%/50%   1         10        1          14m
php-apache    Deployment/php-apache     250%/50%   1         10        1          14m
php-apache    Deployment/php-apache     250%/50%   1         10        4          14m
php-apache    Deployment/php-apache     158%/50%   1         10        5          14m
php-apache    Deployment/php-apache     69%/50%   1         10        5          15m
^Chem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl get deployment php-apache
NAME          READY  UP-TO-DATE  AVAILABLE  AGE
php-apache    5/5    5           5           4m42s
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl get hpa php-apache --watch
NAME          REFERENCE                TARGETS  MINPODS  MAXPODS  REPLICAS  AGE
php-apache    Deployment/php-apache     0%/50%   1         10        5          16m
^Chem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl get deployment php-apache
NAME          READY  UP-TO-DATE  AVAILABLE  AGE
php-apache    5/5    5           5           5m44s
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl get deployment php-apache
NAME          READY  UP-TO-DATE  AVAILABLE  AGE
php-apache    5/5    5           5           6m21s
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl get hpa php-apache --watch
NAME          REFERENCE                TARGETS  MINPODS  MAXPODS  REPLICAS  AGE
php-apache    Deployment/php-apache     0%/50%   1         10        5          18m
^Chem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl get deployment php-apache
NAME          READY  UP-TO-DATE  AVAILABLE  AGE
php-apache    5/5    5           5           7m39s
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl get deployment php-apache
NAME          READY  UP-TO-DATE  AVAILABLE  AGE
php-apache    5/5    5           5           8m15s
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl get hpa php-apache --watch
NAME          REFERENCE                TARGETS  MINPODS  MAXPODS  REPLICAS  AGE
php-apache    Deployment/php-apache     0%/50%   1         10        1          21m
^Chem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$ kubectl get deployment php-apache
NAME          READY  UP-TO-DATE  AVAILABLE  AGE
php-apache    1/1    1           1           10m
hem_distributedsystems@cloudshell:~ (cohesive-cell-362020)$
```

There are 3 types of autoscaling features in Kubernetes:

- i) Horizontal Pod Autoscaler
- ii) Vertical Pod Autoscaler
- iii) Cluster Autoscaler

For the Design section of this lab, I chose to work on Horizontal Pod Autoscaler (HPA). HPA helps with the automation of workload and demand management. That is, it automatically updates the deployment or the statefulset (workload resources) where increasing the load would lead to the deployment of more pods. With a subsequent decrease in load, HPA will ask the workload resource to scale back down as the number of Pods available is way higher than the minimum number specified in the configuration. The screenshot above shows how the number of replicas (pod deployment) increase and decrease according to the load. In the screenshot the load is indicated by the 'Targets' value.

The difference between auto scaling and load balancing is that an Auto Scaler allows automatic scaling up and down while a Loadbalancer distributes the incoming traffic across multiple targets. Meanwhile a request splitter is more of a controller for splitting incoming requests/traffic to different targets.

I have used the 2 following links to complete the Design section:

- i) <https://kubernetes.io/docs/tasks/run-application/horizontal-pod-autoscale/>
- ii) <https://kubernetes.io/docs/tasks/run-application/horizontal-pod-autoscale-walkthrough/>

Video Links:

https://drive.google.com/drive/folders/1gyiVFTrn_n7ga4bEZA1sKPq0ZwX_5YH?usp=sharing